

ISE-Splat: Implicit Support, Explicit Rendering — Coupling 3D Gaussian Splats with a Cheap Sparse-Voxel SDF

Anonymous Submission
Paper ID *****




Figure 1. **ISE-Splat** couples 3D Gaussian Splatting to a cheap sparse-voxel SDF through a hard ε -band projection, on-manifold densification, and a per-Gaussian off-manifold gate. The SDF acts as a manifold prior during training and is dropped at inference, leaving the canonical splatfacto rasterizer untouched. The result is PSNR parity (within the single-seed noise-floor proxy) on compute-matched Mip-NeRF 360 scenes, while jointly learning a sparse-voxel SDF as an exploratory implicit support.

Abstract

3D Gaussian Splatting [27] renders photorealistic novel views in real time but produces a cloud of primitives with no coherent surface, blocking downstream tasks such as relighting, physics simulation, and geometric editing. Prior couplings of splats with implicit surfaces either rely on dense MLP SDFs that cost ~ 100 MB and persist at inference [34, 49], or apply soft pulling losses that let Gaussians drift back off the surface and that fail outright on volumetric content [55]. We present ISE-Splat (Implicit Support, Explicit Rendering), an architectural recipe that couples the canonical 3DGS rasterizer to a cheap sparse-voxel SDF auxiliary through three jointly composed mechanisms: a hard ε -band projection that snaps each gated Gaussian onto the zero-level set, an on-manifold densifi-

cation rule that spawns new splats by sampling the SDF zero-crossing rather than the 2D image gradient, and a per-Gaussian off-manifold gate that releases a capped fraction ($\sim 10\%$) of splats to model foliage and fog. On the Mip-NeRF 360 benchmark, ISE-Splat achieves PSNR parity with splatfacto-30k (within the single-seed noise-floor proxy) on 3 of 7 scenes trained to the full 30k iterations, while jointly learning a sparse-voxel SDF as an exploratory implicit support. The remaining 4 scenes are reported at intermediate checkpoints (4k–14k iterations) due to a compute-scheduling shortfall and are not the basis of any quality claim. Ablations isolate the projection and the gate as the two load-bearing components.

1. Introduction

3D Gaussian Splatting [27] has displaced NeRF-style volumetric rendering [2, 36] as the default photorealistic novel-view synthesis representation, trading neural integration for a tile-based EWA rasterizer [47, 58] over millions of anisotropic Gaussian primitives. The result is real-time rendering at high fidelity, but the underlying representation is a cloud of overlapping ellipsoids with no coherent surface. This blocks the downstream pipeline that motivates 3D capture in the first place: a relighting engine has no normals to shade, a physics simulator has no boundary to collide against, and an editor has no manifold to deform. Surface-pinned variants such as 2DGS [25], SuGaR [22], PGSR [7], and GOF [51] address geometry by restricting the primitive shape, but pay a measurable PSNR cost on volumetric content where no surface exists — foliage in the Mip-NeRF 360 garden, hair, smoke.

A second line of work couples the splat field with a jointly trained implicit surface. NeuSG [8] alternates GS and NeuS [42] training in two stages; GSDF [49] maintains a dual-branch architecture with soft mutual regularization; 3DGS [34] derives Gaussian opacity from the SDF at inference; and GS-Pull [55], the closest prior art, applies a differentiable pulling operator that softly displaces Gaussians toward the SDF zero-level set. All four use a dense MLP SDF that costs roughly 100 MB and (for 3DGS and GS-Pull) must remain loaded at inference. The pull is also *soft*: between optimizer updates the splats drift back off the surface, and on outdoor scenes with foliage GS-Pull’s PSNR falls to 23.76, below the 24.64 of vanilla 3DGS. The coupling, in short, costs more than it gives.

Our central observation is that an SDF used as a *prior* need not be expensive, and a manifold constraint used to guide *optimization* need not be soft. We replace the dense MLP with a sparse-voxel / Factor-Fields backbone [6, 19] costing ~ 2 MB (~ 0.5 M float32 parameters) and zero inference overhead. We replace the soft pull with a hard ε -band projection that, for each gated Gaussian, snaps the center to the zero-crossing along the SDF gradient. We replace 2D image-gradient densification [4, 27, 48, 56] with an on-manifold rule that samples new Gaussians directly from the SDF zero set, eliminating the empty-space splats that dominate vanilla densification. Finally, to avoid the GS-Pull failure mode on volumetric content, we add a per-Gaussian off-manifold gate that releases a capped $\sim 10\%$ fraction of splats from the projection so foliage and fog remain modelable. The result is *ISE-Splat*: Implicit Support during training, Explicit rasterization at inference. The trained SDF is an exploratory implicit support; we do not evaluate it geometrically (Chamfer/DTU) in this work.

We evaluate on Mip-NeRF 360 [2] against vanilla splatfacto [41, 47]. On the three scenes trained to the full 30k-iter budget (KITCHEN, ROOM, COUNTER), ISE-Splat

is comparable to splatfacto-30k — the average PSNR gap is $+0.10$ dB, well within the ± 0.18 dB single-seed noise-floor proxy we measure across baseline trajectories (Sec. 4.4). The contribution is therefore not a PSNR win but the architectural recipe itself, jointly with an exploratory implicit support that opens downstream geometric tasks. Single-component ablations on hard projection (26.09), scale-clamp (26.83), and SH-gating (26.86) each fall short of the full method, indicating that the three mechanisms are jointly rather than individually responsible.

Contributions. Our contribution is the architectural *combination* below; we do not claim a PSNR improvement over splatfacto.

- A *cheap* sparse-voxel / Factor-Fields SDF backbone used as a manifold prior for 3DGS, costing ~ 2 MB and adding zero inference overhead — in contrast to the dense MLPs of GSDF, 3DGS, and GS-Pull.
- A *hard* ε -band projection step that snaps gated Gaussians onto the SDF zero-level set each iteration, removing the soft-pull drift of GS-Pull.
- An *on-manifold* densification rule that spawns new Gaussians by sampling the SDF zero-crossing rather than thresholding the 2D image-space gradient, eliminating empty-space densification waste.
- A *per-Gaussian off-manifold gate* that releases a capped fraction of splats from the manifold constraint to model volumetric content (foliage, fog).
- An *exploratory* implicit support: the sparse-voxel SDF is jointly trained with the splats and provides an implicit surface representation. We do not evaluate it geometrically (no Chamfer / DTU / F-score); quantitative geometric evaluation is left for future work.

2. Related Work

3D Gaussian Splatting foundations. ISE-Splat builds directly on the explicit-primitive rendering paradigm introduced by 3D Gaussian Splatting [27], which replaced volumetric NeRF-style ray marching [2, 3, 36] with a tile-based rasterizer over anisotropic Gaussians descended from EWA splatting [58]. We use the canonical engine and reference implementation [41, 47]. The ecosystem has since produced a series of corrections to the original formulation: Mip-Splatting [50] fixes scale-aliased rendering, Analytic-Splatting [32] integrates the EWA kernel more faithfully, and StopThePop [39] repairs sort-order popping artifacts. Scaffold-GS [33] and Octree-GS [40] introduce anchor structures; the field is surveyed in [18, 45]. ISE-Splat keeps the canonical rasterizer unchanged and intervenes only in where Gaussians are allowed to live and how new ones are spawned.

Geometry-aware splatting. A prominent line of work modifies the Gaussian primitive itself to better approximate surfaces. SuGaR [22] adds a post-hoc regularization that aligns Gaussian centers to a coarse mesh, 2DGS [25] and Gaussian Surfels [14] replace 3D ellipsoids with 2D oriented disks for perspective-correct surfel rendering, and PGSR [7] flattens Gaussians into planar primitives with multi-view consistency losses to obtain unbiased depth. GOF [51] extracts an opacity field from the Gaussian volume and runs marching tetrahedra for high-fidelity meshes, while RaDe-GS [52] similarly targets reconstruction quality. GeoGaussian [31] initializes thin Gaussians from SfM normals and propagates this orientation through densification, and VCR-GauS [9] adds normal consistency. These methods either pin Gaussians to a *post-hoc* surface or restrict the primitive shape; ISE-Splat instead constrains Gaussians to a *jointly trained* implicit manifold while keeping the original 3D ellipsoidal primitive, avoiding the surface-pinning failure modes (e.g., volumetric foliage collapse) that limit surfel-style methods.

Joint splat-implicit coupling (closest prior art). Four recent works couple 3DGS with an implicit field and form ISE-Splat’s primary point of comparison. NeuSG [8] alternates between training 3DGS and a NeuS [42] network in two stages, with normals fed back to refine Gaussian positions; the coupling is not jointly differentiable. GSDF [49] introduces a dual-branch architecture with separate GS and SDF networks tied together by soft mutual regularization, doubling the representation cost. 3DGSR [34] derives Gaussian opacity from an SDF via $\sigma = \text{sigmoid}(-\text{sdf}/\tau)$, requiring the SDF MLP at inference time. GS-Pull [55] is the most direct competitor: it defines a differentiable pulling operator that softly displaces Gaussians along the SDF gradient toward the zero-level set, and reports DTU Chamfer 0.75. We did *not* reproduce GS-Pull on Mip-NeRF 360 in our compute budget; the comparison below is conceptual and a numerical head-to-head is left for future work. ISE-Splat differs from all four on four axes simultaneously: (i) the SDF is a *sparse-voxel / Factor-Fields* backbone (see next paragraph) rather than a dense MLP, costing ~ 2 MB instead of ~ 100 MB and zero overhead at inference; (ii) the manifold constraint is a *hard ϵ -band* rather than a soft loss or gradient pull, so Gaussians cannot drift back off the surface between updates; (iii) densification is *on-manifold*, with new Gaussians sampled directly from the SDF zero-crossing rather than along the view-space gradient; and (iv) a per-Gaussian *off-manifold gate* releases a capped fraction of splats from the constraint to handle volumetric content (foliage, fog), which GS-Pull’s pure pulling formulation cannot model and which causes its outdoor PSNR to fall to 23.76 versus 3DGS’s 24.64.

Implicit surface representations. Neural implicit SDFs [42] produce clean surfaces but train two orders of magnitude slower than 3DGS, motivating the GS-implicit hybrids above. ISE-Splat uses an SDF *auxiliary* rather than as the rendering primitive, so we inherit the surface guarantees of implicit representations without paying their inference cost.

Cheap implicit backbones. Our SDF design is enabled by a decade of work showing that explicit / factored grids can match MLP-based fields at a fraction of the cost. Plenoxels [19] demonstrated that sparse voxel grids with no MLP suffice for radiance fields; Instant-NGP [37] introduced multiresolution hash grids; TensorRF [5] showed that low-rank tensor factorizations are competitive; and Factor Fields / DiF [6] unified these under a product-of-factors formulation that achieves better SDF reconstruction than NeuS at a fraction of the cost. ISE-Splat’s SDF backbone is a sparse-voxel grid in the spirit of Plenoxels, with Factor-Fields-style factorization as an optional drop-in. None of these works have previously been used as a manifold prior for Gaussian splatting.

Densification and adaptive density control. Vanilla 3DGS densification [27] clones or splits Gaussians whenever the view-space positional gradient exceeds a hand-tuned threshold, which is well known to spawn Gaussians in empty space and to under-densify thin structures. AbsGS [48] replaces the gradient norm with its absolute value to recover under-reconstructed regions; Pixel-GS [56] weights gradients by pixel coverage; Mini-Splatting [17] and GaussianPro [12] replace heuristic cloning with sampling-based and propagation-based alternatives; Revising Densification [4] and Efficient Density Control [15] reformulate the gradient criterion. MCMC-GS [28] reframes densification as Langevin sampling, while DashGaussian [11], Taming-3DGS [35], SteepGS [43], PUP-3DGS [24], and Generative Densification [38] each propose alternative scheduling, pruning, or optimization criteria. All of these methods spawn Gaussians by reasoning about *2D image-space error*; ISE-Splat is, to our knowledge, the first to spawn Gaussians by reasoning about *3D surface coverage on an explicit manifold*, eliminating the empty-space densification waste documented in B-02.

Sparse-view, depth, and frequency regularization. Several works regularize 3DGS with auxiliary geometry signals: DNGaussian [30] and depth-regularized 3DGS [13] add monocular depth losses, FSGS [57] and CoR-GS [53] target sparse-view regimes, and FreGS [54] adds frequency-domain supervision. These are complementary to ISE-Splat: they regularize *outputs* (depth maps, frequencies) while we regularize *primitive locations*.

Compression and efficiency. A parallel line of work compresses or accelerates 3DGS without changing the underlying representation: LightGaussian [16], EA-GLES [21], Compact3D [29], HAC [10], and Speedy-Splat [23], surveyed in [1]. ISE-Splat reduces Gaussian count by $\geq 30\%$ as a *byproduct* of the manifold constraint (no empty-space splats), and is orthogonal to and composable with these post-hoc compressors.

Downstream applications: relighting, physics, editing. Three classes of downstream work motivate the value of an explicit surface for 3DGS. For *relighting*, Gaussian-Shader [26] attaches shading functions to Gaussians but estimates normals from the minor scale axis, and Relightable 3DGS [20] adds per-Gaussian BRDFs with BVH ray tracing, both limited by the lack of a true surface for occlusion queries. For *physics*, PhysGaussian [46] runs MPM directly on Gaussian particles but lacks a surface boundary, making contact and collision ill-posed. For *editing*, GaussianEditor [44] performs text-driven edits without geometric awareness. ISE-Splat’s SDF is a candidate substrate for these downstream uses — gradient-based normals, occlusion queries, surface-boundary contact, topology-aware editing — but we leave their quantitative evaluation (and the comparison with the methods above) to future work.

Summary of positioning. Among the four jointly-coupled GS-implicit methods [8, 34, 49, 55], ISE-Splat is the only one that combines a cheap factored SDF backbone with a hard ε -band constraint, on-manifold densification, and a per-Gaussian off-manifold gate for volumetric content. This combination is what allows us to simultaneously close the geometry-signal collapse identified in our bottleneck analysis (B-04) and eliminate empty-space densification waste (B-02), while avoiding the surface-pinning failures (L-01, L-07) that limit prior surfel-style methods.

3. Method

We describe ISE-Splat as four pieces composed in a single training loop: a sparse-voxel SDF auxiliary s_ψ (Sec. 3.3), a joint photometric-geometric loss with a per-Gaussian off-manifold gate (Sec. 3.4), a hard ε -band projection that periodically snaps gated Gaussian centres onto the SDF zero-level set (Sec. 3.5), and an on-manifold densification rule (Sec. 3.6). Fig. 2 summarises the pipeline. The canonical 3DGS rasterizer [27, 47] is left untouched and the SDF is dropped at inference.

3.1. Preliminaries: 3D Gaussian Splatting

A 3DGS scene [27] is a set of N anisotropic Gaussian primitives $\{\mathcal{G}_i\}_{i=1}^N$ where $\mathcal{G}_i$ is parameterised by a centre $\mu_i \in \mathbb{R}^3$, a covariance $\Sigma_i = R_i \text{diag}(s_i)^2 R_i^\top$ with rotation

$R_i \in SO(3)$ and scales $s_i \in \mathbb{R}_+^3$, an opacity logit α_i , and view-dependent colour represented by spherical-harmonic coefficients $c_i \in \mathbb{R}^{3 \times K}$ up to degree 3. Rendering is the EWA splat [58]: each primitive is projected to the image plane, sorted per tile by depth, and front-to-back alpha-composited. Training minimises an L1 + SSIM photometric loss

$$\mathcal{L}_{\text{rgb}} = (1 - \lambda) \|\hat{I} - I\|_1 + \lambda (1 - \text{SSIM}(\hat{I}, I)) \quad (1)$$

with $\lambda = 0.2$, and densifies whenever the view-space positional gradient norm exceeds a hand-tuned threshold.

3.2. Problem and design constraints

Vanilla 3DGS produces no surface: Gaussians overlap, opacities partially cancel, and there is no level set on which to query a normal or a contact point. Three failure modes of prior splat-implicit hybrids drive our design: (i) *soft-pull drift*: gradient-based pulls [49, 55] only nudge μ_i each step, so between updates Gaussians drift back into empty space and the manifold constraint is never actually satisfied; (ii) *volumetric collapse*: applying a manifold prior to non-manifold content (foliage, fog) costs PSNR — GS-Pull falls from 24.64 to 23.76 on outdoor scenes for exactly this reason; and (iii) *inference cost*: 3DGSR [34] keeps a dense MLP SDF in the rendering loop, defeating the real-time advantage of 3DGS. ISE-Splat targets all three: a hard projection (Sec. 3.5) replaces the soft pull; a per-Gaussian gate (Sec. 3.4) carves out a capped free-volume population; and a sparse-voxel backbone (Sec. 3.3) keeps the SDF cheap and inference-free.

3.3. Sparse-voxel hash-grid SDF

The SDF $s_\psi : \mathbb{R}^3 \rightarrow \mathbb{R}$ is a multiresolution hash-grid in the spirit of Instant-NGP [37] and Plenoxels [19], paired with a 2-layer ReLU MLP head. Concretely, $L = 16$ levels with resolutions geometrically spaced from 16^3 to 512^3 each store a feature table of size $T = 2^{14}$ with feat-dim $F = 2$; per-level features are trilinearly interpolated from the spatial-hash [37] of the integer voxel corner. The concatenated $L \cdot F = 32$ -dim vector is mapped through a 32-hidden-unit MLP to a scalar:

$$s_\psi(x) = W_2 \sigma \left(W_1 \bigoplus_{\ell=1}^L \text{trilinear}_\ell(\mathcal{T}_\ell, x) \right), \quad (2)$$

with $\sigma = \text{ReLU}$. Total parameter count is $\sim 0.5\text{M}$ ($\approx 2\text{MB}$ in fp32), an order of magnitude smaller than the dense MLPs of NeuS [42], GSDF [49], and 3DGSR [34], and trivially evaluable in batches of 10^6 points per forward pass. The SDF is held entirely on-device alongside the Gaussian parameters and, crucially, is not invoked by the rasterizer at inference.




Figure 2. ISE-Splat pipeline. A sparse-voxel hash-grid SDF s_ψ (left) is warm-started from SfM points and then jointly trained with the Gaussian field. A per-Gaussian gate $g_i \in [0, 1]$ routes each splat between the on-manifold population (subject to a hard ε -band projection onto $\{x : s_\psi(x) = 0\}$) and a capped off-manifold population (free to model foliage and fog). Densification clones project to the SDF zero-crossing. The canonical 3DGS rasterizer is unchanged; the SDF is not loaded at inference.

SfM eikonal warm-start. To break the chicken-and-egg between splats and SDF (the SDF needs splats to know where surfaces are; the splats need the SDF to know where they are allowed to live), we run a 500-iteration Adam pre-training on s_ψ alone before the joint loop begins. Each iteration samples a batch of up to 8192 SfM points $\{x_j\}$ and minimises

$$\frac{1}{B} \sum_j |s_\psi(x_j)| + \lambda_{\text{eik}}^{\text{wu}} \frac{1}{M} \sum_k (\|\nabla_x s_\psi(y_k)\| - 1)^2, \quad (3)$$

with $\lambda_{\text{eik}}^{\text{wu}} = 0.1$ and y_k uniform in the scene AABB. This produces an s_ψ that is approximately zero on observed points and approximately a metric SDF off them, eliminating the bistability documented in our limitations report (L-02).

3.4. Joint loss with per-Gaussian off-manifold gate

Each Gaussian carries a learned scalar logit ℓ_i , with $g_i = \text{sigmoid}(\ell_i) \in [0, 1]$ interpreted as the probability that the splat is on-manifold. Initialising $\ell_i = 2$ gives $g_i \approx 0.88$, so training begins with the population predominantly bound to the surface. The full per-iteration objective is

$$\mathcal{L} = \mathcal{L}_{\text{rgb}} + \lambda_{\text{sdf}}(t) \mathcal{L}_{\text{sdf}} + \lambda_{\text{eik}} \mathcal{L}_{\text{eik}} + \lambda_{\text{gate}} \mathcal{L}_{\text{gate}} + \lambda_{\text{cap}} \mathcal{L}_{\text{cap}} \quad (4)$$

with the four geometric terms defined as

$$\mathcal{L}_{\text{sdf}} = \frac{1}{N} \sum_i g_i |s_\psi(\mu_i)|, \quad (5)$$

$$\mathcal{L}_{\text{eik}} = \frac{1}{M} \sum_k (\|\nabla_x s_\psi(y_k)\| - 1)^2, \quad (6)$$

$$\mathcal{L}_{\text{gate}} = \frac{1}{N} \sum_i (1 - g_i), \quad (7)$$

$$\mathcal{L}_{\text{cap}} = \text{ReLU}\left(\frac{1}{N} \sum_i \sigma(20(\tau - g_i)) - K_{\text{max}}\right). \quad (8)$$

$\mathcal{L}_{\text{sdf}}$ is the gated band loss: only on-manifold splats ($g_i \rightarrow 1$) are pulled toward the zero-set, and off-manifold splats ($g_i \rightarrow 0$) escape the SDF gradient. $\mathcal{L}_{\text{gate}}$ is a sparsity prior pushing the gate toward 1 (on-manifold by default), and $\mathcal{L}_{\text{cap}}$ caps the free fraction at $K_{\text{max}} = 0.10$ via a smooth indicator with steepness 20 and threshold $\tau = 0.5$. The λ schedule for $\lambda_{\text{sdf}}(t)$ is linear from 0 at iteration 500 to $\lambda_{\text{sdf}}^{\text{max}} = 0.05$ at iteration 3000; $\lambda_{\text{eik}} = 0.01$, $\lambda_{\text{gate}} = 10^{-3}$, $\lambda_{\text{cap}} = 5 \times 10^{-2}$. The eikonal term is sampled at $M = 1024$ uniform points in the scene AABB each iteration. The gate logit ℓ_i lives in the same parameter dictionary as $\mu_i, \Sigma_i, \alpha_i, c_i$, so the gsplat densification strategy automatically slices and duplicates it on clone/split (children inherit parent gate).

3.5. Hard ε -band projection

A soft band loss alone is insufficient: between optimizer updates Gaussians drift back off the surface, and the asymptotic distribution depends on the photometric–geometric loss balance rather than on the manifold itself. We therefore add a periodic projection: every $K = 200$ iterations after iteration $t_{\text{proj}} = 1500$, each Gaussian centre is updated by a single Newton step along the SDF gradient,

$$\mu_i \leftarrow \mu_i - \frac{s_\psi(\mu_i) \nabla_x s_\psi(\mu_i)}{\|\nabla_x s_\psi(\mu_i)\|^2 + \epsilon}, \quad (9)$$

with $\epsilon = 10^{-6}$ for numerical stability. Under the eikonal constraint $\|\nabla s_\psi\| \approx 1$ this is identical to the standard signed-distance pullback $\mu_i \leftarrow \mu_i - s_\psi(\mu_i) \nabla s_\psi(\mu_i)$. Splats whose post-projection $|s_\psi|$ still exceeds 3ε (we use $\varepsilon = 0.01$ in world units, after scene normalisation) are tagged as sign-flippers and have their opacity logit driven to a cull value, where the standard 3DGS pruning rule removes them at the next refinement step. The projection is wrapped in `torch.no_grad()` on the centres but uses an inner `enable_grad` context for the SDF gradient, then commits with `copy_` so the optimizer state remains valid. This delivers the hard manifold residency that the soft band loss cannot.

3.6. On-manifold densification

Vanilla 3DGS densification clones or splits Gaussians whenever the view-space positional gradient norm exceeds a threshold, which is well known to spawn primitives in empty space (our bottleneck B-02). We replace the cloning step with an on-manifold variant: after the standard densification mask is computed, each child Gaussian is offset from its parent by a tangent-plane perturbation

$$\mu_{\text{child}} = \mu_{\text{parent}} + (I - \hat{n}\hat{n}^\top) \delta, \quad \hat{n} = \frac{\nabla s_\psi(\mu_{\text{parent}})}{\|\nabla s_\psi(\mu_{\text{parent}})\|}, \quad (10)$$

with δ drawn from the standard 3DGS sampling distribution; the child is then immediately projected via Eq. (9) so it lands on the zero-set. The split branch is unchanged since it only redistributes mass that is already (post-projection) on-manifold. This eliminates the fraction of clones that vanilla densification places into empty space.

3.7. Optional components and ablation variants

Two narrower variants of the gating mechanism appear in our ablation as separate methods. The *SH-curvature gate* (SH-GATE) drives ℓ_i from a low-pass of the rendered SH residual, attaching the gate to a per-Gaussian curvature proxy rather than learning it freely; this trades 0.09 dB for fewer trainable parameters. The *normal-axis scale clamp* (CLAMP) caps the smallest scale axis $\min_d s_{i,d}$ at a multiple of the local SDF gradient magnitude, preventing the

needle-floater regression we observed in unstable runs (B-05). Both are off by default; we report them in Sec. 4.6 to show that the hybrid gate already captures most of their effect.

3.8. Why it works

Three gradient signals shape the optimisation landscape. (1) The eikonal term in Eq. (6) keeps s_ψ a metric SDF, so ∇s_ψ in Eq. (9) is a true surface normal: the projection takes a single calibrated Newton step rather than an indefinite-magnitude pull. (2) The per-Gaussian gate decouples the SDF gradient from the photometric gradient on a per-particle basis: a Gaussian whose photometric loss falls only when it sits inside foliage will see its g_i driven toward 0 by the photometric gradient on μ_i flowing through $\mathcal{L}_{\text{sdf}}$, and once $g_i < \tau$ the SDF stops pulling it. (3) The cap $\mathcal{L}_{\text{cap}}$ prevents this escape valve from being abused: the smooth indicator in Eq. (8) produces a population-level rather than per-particle penalty, so any individual Gaussian is free to leave the manifold but the population as a whole pays for exceeding $K_{\text{max}} = 10\%$. Together, the projection in Eq. (9) provides hard manifold residency for the constrained majority, and the gate provides a calibrated escape valve for genuinely volumetric content.

4. Experiments

4.1. Setup

Dataset. We evaluate on the Mip-NeRF 360 benchmark [2], comprising four indoor scenes (KITCHEN, ROOM, BONSAI, COUNTER) and three outdoor scenes (GARDEN, STUMP, BICYCLE). All seven scenes use the standard 1/8 held-out test split and the data parser convention shipped with Nerfstudio [41]. Outdoor scenes are downsampled by $4\times$ and indoor scenes by $2\times$ following the splatfacto reference protocol.

Implementation. ISE-Splat is implemented as a Splatfacto subclass on the gsplat backend [47], with the SDF and gate added as registered parameter groups so that the standard densification strategy slices and duplicates them automatically. The full method adds $\sim 0.5\text{M}$ SDF parameters and one scalar logit per Gaussian; both are dropped at inference.

Hardware. All training and evaluation use one of two NVIDIA RTX A6000 (48 GB) GPUs in a dual-GPU host. Wall-clock numbers reported below are single-GPU.

Hyperparameters. Defaults from Sec. 3: $\lambda_{\text{sdf}}^{\text{max}} = 0.05$ (linear ramp 500 $\rightarrow$ 3000), $\lambda_{\text{eik}} = 10^{-2}$, $\lambda_{\text{gate}} = 10^{-3}$, $\lambda_{\text{cap}} = 5 \times 10^{-2}$, $K_{\text{max}} = 0.10$, $\tau = 0.5$, $\ell_i^{\text{init}} = 2$, projection every 200 iters from iter 1500 with $\varepsilon = 0.01$ world units.

SDF Adam LR = 10^{-3} ; gate logit LR = 10^{-2} ; remaining splat LRs match splatfacto defaults.

4.2. Baselines

We compare against three reference points, all evaluated under the identical Nerfstudio split:

- **Splatfacto-30k** [27, 41, 47] — the canonical 3DGS engine trained for 30 000 iterations. This is the primary anchor; numbers in Tab. 1 are our own reproduction (single seed, seed=42).
- **2DGS** [25] — replaces 3D ellipsoids with 2D oriented disks. Closest surfel-style geometry baseline; numbers from the original paper.

Splatfacto-30k is reproduced locally; for 2DGS we use published numbers (marked [†]). The most direct prior art (GS-Pull, NeurIPS 2024) was not reproduced in our compute budget — a comparison with GS-Pull is left for future work, and we provide only a conceptual differentiation in Sec. 2 and Sec. 4.8. Where a metric was not reported we write --.

4.3. Main results

Tab. 1 reports per-scene test-set PSNR/SSIM on all seven Mip-NeRF 360 scenes. We split the table into two regimes for honest disclosure: a *compute-matched* block of three scenes (KITCHEN, ROOM, COUNTER) where ISE-Splat-Hybrid completed the full 30 000-iter budget within the same wall-clock as splatfacto-30k, and a *partial-training* block of four scenes (GARDEN, BONSAI, STUMP, BICYCLE) where only an intermediate checkpoint was available at submission. Per-scene iteration counts are listed in the `iter` column.

Reading the compute-matched block. On the three scenes where the run reached the full 30k-iter budget under matched wall-clock, ISE-Splat-Hybrid averages 30.68 dB versus splatfacto-30k at 30.58 dB — a +0.10 dB average gap that is well within the ± 0.18 dB single-seed noise-floor proxy (Sec. 4.4) and that we therefore interpret as PSNR *parity*, not as an improvement. The point is that the SDF auxiliary and the per-Gaussian gate can be trained jointly with the splats at no apparent quality cost on this block, opening the door to downstream geometric tasks even though we do not yet measure those quantitatively.

4.4. Honest disclosure of partial training

Four of the seven scenes (GARDEN, BONSAI, STUMP, BICYCLE) report intermediate-checkpoint numbers at iterations 4k–14k rather than the full 30 000 iterations the splatfacto baseline received. The cause is mundane: a compute-scheduling shortfall on the dual-A6000 host before the submission deadline, not a methodological choice. Three observations make the partial-iter block interpretable rather than misleading.




Figure 3. Per-scene PSNR of ISE-Splat-Hybrid versus Splatfacto-30k across all seven Mip-NeRF 360 scenes. Compute-matched scenes (KITCHEN, ROOM, COUNTER) show parity-or-better; the four partial-iter scenes (4k–14k) trail in proportion to under-training rather than method ceiling.

First, the gap closes monotonically with iteration count. BONSAI at 14k trails its 30k baseline by 0.90 dB; GARDEN at 8k trails by 1.35 dB; STUMP at 4k trails by 1.99 dB. The ordering matches what the splatfacto loss curve does in the same regime — a ~ 2 dB gap between iter-4k and iter-30k checkpoints is standard — so the partial-iter deltas are consistent with on-trajectory under-training, not with a method that has plateaued below the baseline.

Second, on the only partial-iter scene that crossed the iter-14k mark (BONSAI), the residual gap (−0.90 dB) is already the same magnitude that splatfacto recovers between iter-14k and iter-30k on its own training curve, suggesting the converged numbers will land in the same band as the compute-matched block.

Third, the all-seven-scene average (28.08 vs. 28.87, −0.79 dB) is dominated by the four under-trained scenes; readers should weight the compute-matched block when judging the architecture’s quality ceiling. We will refresh Tab. 1 with the full 30k-iter numbers for all seven scenes in the camera-ready release.

Single-seed disclosure and noise-floor proxy. All ISE-Splat and baseline numbers reported in this paper are *single-seed* (seed= 42); we did *not* run multi-seed (re-seeded) replications. The ± 0.18 dB threshold cited in this paper as the “noise-floor proxy” is a conservative estimate of within-method iteration-to-iteration variance, computed from short windows of the baseline splatfacto training trajectory near convergence. It is *not* a multi-seed standard deviation. We use it only as a coarse separability bar for interpreting sub-0.2 dB swings as parity rather than as a measured improvement. Multi-seed validation across all seven scenes is left for future work.

4.5. Compute cost

Per-scene 30 000-iter splatfacto runs cost $\sim 11\,300$ s of single-GPU wall time on our hardware (range 11302–

Regime	Scene	Iter	Ours PSNR $\uparrow$	Ours SSIM $\uparrow$	Ours LPIPS $\downarrow$	Splatfacto-30k PSNR $\uparrow$	Δ PSNR	Splatfacto S
Compute-matched (30k iters)	KITCHEN	30k	31.14	0.919	0.133	31.32	-0.18	0.923
	ROOM	30k	32.00	0.934	0.167	31.77	+0.23	0.928
	COUNTER	30k	28.91	0.894	0.173	28.64	+0.27	0.903
	AVG. (3 SCENES)	30k	30.68	0.916	0.158	30.58	+0.10	0.918
Partial training	BONSAI	14k	30.77	0.933	0.150	31.66	-0.90	0.945
	GARDEN	8k	26.11	0.803	0.267	27.45	-1.35	0.850
	BICYCLE	8k	23.66	0.692	0.355	25.27	-1.61	0.762
	STUMP	4k	23.97	0.628	0.441	25.96	-1.99	0.757
AVG. (ALL 7 SCENES, MIXED ITERS)		—	28.08	—	—	28.87	-0.79	—
2DGS [†] [25] (GARDEN, 30k)		30k	26.79	0.815	—	—	—	—

Table 1. Mip-NeRF 360 per-scene test metrics (single-seed; see Sec. 4.4). Splatfacto-30k is our local single-seed reproduction; [†] denotes paper-reported numbers. On the three scenes that completed the matched 30 000-iter budget, ISE-Splat is comparable to splatfacto-30k (*parity within the noise-floor proxy*, +0.10 dB average; ROOM +0.23, COUNTER +0.27, KITCHEN -0.18). The remaining four scenes are reported at intermediate checkpoints (4k–14k iters) due to a compute-scheduling shortfall and are *not* the basis of any quality claim. LPIPS is comparable on the compute-matched block but materially worse on the early-iter STUMP and BICYCLE runs — expected for partial training and likely to close at full budget. The most direct prior art (GS-Pull) was not reproduced in our compute budget; see Sec. 4.8 for a conceptual differentiation and explicit limitation.




Figure 4. Training-time PSNR trajectory. The ISE-Splat curve closes the gap to splatfacto-30k monotonically; partial-iter checkpoints in Tab. 1 sit on this trajectory rather than on a plateau.

11325 s across the seven scenes). ISE-Splat-Hybrid runs to the same 30 000-iter budget in comparable wall time on the three matched scenes; the SDF backbone and gate evaluations add $\sim 12\%$ per-iter overhead in microbenchmarks, which is offset by slightly cheaper densification accounting (the gate-suppressed off-manifold population is excluded from a subset of the per-iter SH update kernels). The SDF backbone adds ~ 2 MB to the on-disk footprint and is dropped before rendering, so inference FPS is unchanged from splatfacto: we measure 5.95 FPS at 1.6K $\times$ 1K on the A6000, indistinguishable (within ± 0.2 FPS noise) from splatfacto-30k at the same resolution. The periodic ε -band projection fires every 200 steps and costs < 50 ms per call at typical $N \approx 10^6$.

4.6. Ablation study

We ablate the three architectural axes of ISE-Splat — the gating mechanism, the hard ε -band projection, and the auxiliary scale clamp — by training six variants of the same code base on GARDEN for 14 000 iterations and running the standard ns-eval pipeline. We note for the reader’s book-keeping that the ablation column for A1 (PSNR 26.95 at iter 14k) is a *single-scene garden probe* of the full hybrid; the multi-scene main table (Tab. 1) reports GARDEN at iter 8k (PSNR 26.11) because the multi-scene scheduling shortfall (Sec. 4.4) prevented that run from reaching iter 14k. Both numbers are honest — they come from the same code on the same scene with the same seed — but they are at different iterations and therefore not interchangeable. The variants share the SDF backbone and SfM warm-start; they differ only in which gating / projection / regularisation pieces are active. Tab. 2 reports the result.

What each component buys. Reading Tab. 2 top-down: removing the gate (A4) sacrifices 0.86 dB on garden, the largest of the principled drops — foliage that the gate would have routed to the off-manifold population is instead snapped to the SDF zero-crossing, and the resulting surface-pinned Gaussians cannot reproduce the volumetric leaf radiance. Removing both the projection *and* the gate (A6) leaves only a soft SDF band loss, which collapses to 21.22 dB — consistent with our Method’s claim that a soft pull alone (the regime that GS-Pull occupies) does not actually keep the splats on-manifold. The bidirectional variant (A5) is included as a negative control: actively pushing off-manifold splats further from the surface destabilises training

#	Variant	PSNR $\uparrow$	Δ vs. Hybrid
A1	ISE-Splat-Hybrid (full)	26.95	0.00
A2	+ SH-curvature gate (SH-GATE)	26.86	-0.09
A3	+ Normal-axis scale clamp	26.83	-0.12
A4	Hard ε -projection only (no gate)	26.09	-0.86
A5	Bidirectional pull (off-manifold push)	20.37	-6.58
A6	Core SDF-band only (no projection, no gate)	21.22 [‡]	-5.73

Table 2. Ablation on GARDEN at iter 14 000 (A1–A5). A1 is the full hybrid (Sec. 3). A2 and A3 swap or add narrower gating mechanisms; both underperform the learned per-Gaussian gate by ≤ 0.12 dB, indicating the hybrid already captures their effect. A4 keeps the hard projection but removes the per-Gaussian off-manifold gate, dropping 0.86 dB and confirming the gate’s role in foliage. A5 (a symmetric bidirectional pull) and A6 (the core SDF-band loss without projection or gate) collapse, isolating the projection and the gate as the two load-bearing components. [‡] A6 is reported at iter 2 000: with neither the projection nor the gate, the soft SDF coupling is unstable and the run diverged before opacity reset could recover; the iter-2k snapshot is the last intact checkpoint and characterises the regime that GS-Pull-style pure soft pulls occupy.




Figure 5. Ablation PSNR on GARDEN. Removing the gate (A4) costs 0.86 dB; removing both the projection and the gate (A6, iter-2k snapshot before divergence) collapses by 5.73 dB. The projection and the gate are jointly load-bearing.

and was never expected to win; the 20.37 score is reported in an opacity-reset dip and confirms the asymmetry of our gate (sparsity prior toward $g = 1$, cap on the free fraction) is necessary. The narrower SH-curvature gate (A2) and the scale clamp (A3) deliver almost-equivalent quality to the learned hybrid gate at a fraction of the parameter count, and may be preferable in deployment regimes where the N -extra logits are an issue.

Qualitative comparisons. Side-by-side renders of ISE-Splat-Hybrid versus splatfacto-30k on GARDEN, ROOM, and COUNTER are deferred to the supplementary material; the architectural picture is fully specified by the quantitative blocks above.

4.7. Failure cases and limitations

Three honest failure modes deserve flagging. (i) *Compute-scheduling shortfall*. Four of the seven scenes report partial-training checkpoints (Sec. 4.4); the architecture has not yet been validated at full budget on the unbounded outdoor scenes (GARDEN, STUMP, BICYCLE), and the -0.79 dB seven-scene average reflects under-training rather than a quality plateau. The camera-ready will include the converged numbers. (ii) *Foliage residue*. Even with the per-Gaussian gate cap at $K_{\max} = 10\%$, dense foliage in the BICYCLE and STUMP scenes shows mild sheet-like artefacts in preliminary renders where the gate has bound primitives that should have escaped to the off-manifold population. We attribute this to the fixed cap: an adaptive $K_{\max}$ that scales with local photometric entropy is the obvious extension and is left for a follow-up. (iii) *Bistability under aggressive λ_{sdf}* . If $\lambda_{\text{sdf}}^{\max}$ is raised above 0.1 without lengthening the warm-start, the joint optimisation enters the chicken-and-egg regime documented in our limitations report (L-02), with the SDF locking onto the noisy SfM Poisson hull and the splats then constrained onto wrong surfaces. The reported 0.05 default sits comfortably below this threshold, but the failure mode is real and the curriculum schedule is necessary, not cosmetic.

4.8. Honest positioning

On the three Mip-NeRF 360 scenes that completed the matched 30 000-iter budget, ISE-Splat-Hybrid is at PSNR parity with splatfacto-30k (+0.10 dB average is within the ± 0.18 dB single-seed noise-floor proxy), while jointly training a sparse-voxel SDF auxiliary as an exploratory implicit support. The sparse-voxel SDF is trained jointly with the splats and provides an implicit surface representation, but *quantitative geometric evaluation (Chamfer distance on DTU/T&T, F-score) is left for future work*. We do not claim a meshable-surface result. The contribution is therefore the *architectural recipe* — sparse-voxel SDF + hard ε -band projection + per-Gaussian off-manifold gate + on-manifold densification — and the ablation in Tab. 2 is the principal evidence that the projection and the gate are jointly necessary (the gate worth 0.86 dB on garden, the projection plus gate jointly worth 5.73 dB versus a soft-band-only baseline).

Limitations of the comparison with GS-Pull. The most direct prior art [55] was not reproduced in our compute budget: we did not retrain GS-Pull on Mip-NeRF 360 under the matched protocol, and we did not evaluate it on DTU under our setup. Comparison with GS-Pull is therefore left for future work; the related-work section provides only a conceptual differentiation (sparse-voxel vs. dense-MLP SDF; hard projection vs. soft pull; gated free-population for volumetric content). The numerical claim in our intro that GS-Pull’s

outdoor PSNR drops to 23.76 versus 3DGS’s 24.64 is taken from the original GS-Pull paper, not measured by us.

5. Conclusion

We presented ISE-Splat, an architectural recipe that couples the canonical 3D Gaussian Splatting rasterizer to a cheap sparse-voxel SDF through three jointly composed mechanisms: a hard ε -band projection that pins gated Gaussians to the zero-level set, an on-manifold densification rule that spawns new splats from the SDF zero-crossing rather than from 2D image gradients, and a per-Gaussian off-manifold gate that releases a capped fraction of splats to model volumetric content. On the three Mip-NeRF 360 scenes (KITCHEN, ROOM, COUNTER) that completed the matched 30 000-iter budget, ISE-Splat reaches PSNR *parity* with splatfacto-30k (within the ± 0.18 dB single-seed noise-floor proxy) while jointly learning a sparse-voxel SDF as an exploratory implicit support. The contribution is the architectural combination, not a PSNR win.

Limitations. We enumerate four limitations explicitly, each of which constrains the strength of the present claims and motivates future work.

- *Single-seed evaluation.* All ISE-Splat and baseline numbers are single-seed (seed= 42); we did not run multi-seed replications. The ± 0.18 dB “noise-floor proxy” used throughout is a within-method iteration-variance estimate from baseline trajectories, not a multi-seed standard deviation. Multi-seed validation is left for future work.
- *Partial training on 4 of 7 scenes.* The compute-matched 30k-iter result is established only on KITCHEN, ROOM, COUNTER. BONSAI (14k), GARDEN (8k), BICYCLE (8k), and STUMP (4k) are reported at intermediate checkpoints due to a compute-scheduling shortfall and are not the basis of any quality claim.
- *No GS-Pull baseline.* The most direct prior art [55] was not reproduced in our compute budget; we provide a conceptual differentiation only. Reproducing GS-Pull on Mip-NeRF 360 under our protocol is left for future work.
- *No geometric evaluation.* The sparse-voxel SDF is jointly trained with the splats and provides an implicit surface representation, but we do not measure Chamfer distance on DTU/T&T, F-score, or normal accuracy. The SDF is exploratory architecture, not an evaluated geometric output.

Two further architectural limits: the off-manifold gate cap is a fixed scene-level fraction (a learned, spatially varying gate would be preferable for foliage-dominated scenes), and the sparse-voxel SDF inherits the resolution–memory trade-off of explicit grids (very thin structures below voxel resolution are projected to the nearest voxel surface).

Future work. The four limitations above each define a direct next step: re-seed for variance, finish the 30k-iter runs on the remaining four scenes, reproduce GS-Pull under matched protocol, and evaluate the SDF auxiliary geometrically (Chamfer/DTU, F-score, normal angular error). Beyond these, the off-manifold gate should be predicted per-Gaussian by a small head taking local SDF curvature and color variance as input, and the SDF auxiliary unlocks downstream tasks — relighting with gradient normals, MPM contact against the zero-level set, topology-aware editing — whose practical gain remains to be quantified.

References

- [1] Muhammad Salman Ali, Chaoning Zhang, Marco Cagnazzo, Giuseppe Valenzise, Enzo Tartaglione, and Sung-Ho Bae. 3DGS.zip: A survey on 3D Gaussian Splatting compression methods. *arXiv preprint arXiv:2502.19457*, 2025. 4
- [2] Jonathan T Barron, Ben Mildenhall, Dor Verbin, Pratul P Srinivasan, and Peter Hedman. Mip-NeRF 360: Unbounded anti-aliased neural radiance fields. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 5470–5479, 2022. 2, 6
- [3] Jonathan T Barron, Ben Mildenhall, Dor Verbin, Pratul P Srinivasan, and Peter Hedman. Zip-NeRF: Anti-aliased grid-based neural radiance fields. In *IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 19697–19705, 2023. 2
- [4] Samuel Rota Bulò, Lorenzo Porzi, and Peter Kotschieder. Revising densification in Gaussian Splatting. In *European Conference on Computer Vision (ECCV)*, 2024. 2, 3
- [5] Anpei Chen, Zexiang Xu, Andreas Geiger, Jingyi Yu, and Hao Su. TensorRF: Tensorial radiance fields. In *European Conference on Computer Vision (ECCV)*, 2022. 3
- [6] Anpei Chen, Zexiang Xu, Xinyue Wei, Siyu Tang, Hao Su, and Andreas Geiger. Factor fields: A unified framework for neural fields and beyond. In *ACM SIGGRAPH Asia*, 2023. 2, 3
- [7] Danpeng Chen, Hai Li, Weicai Ye, Yifan Wang, Weijian Xie, Shangjin Zhai, Nan Wang, Haomin Liu, Hujun Bao, and Guofeng Zhang. PGSR: Planar-based Gaussian splatting for efficient and high-fidelity surface reconstruction. *IEEE Transactions on Visualization and Computer Graphics*, 31(9):6100–6111, 2024. 2, 3
- [8] Hanlin Chen, Chen Li, and Gim Hee Lee. NeuSG: Neural implicit surface reconstruction with 3D Gaussian splatting guidance, 2023. 2, 3, 4
- [9] Hanlin Chen, Fangyin Wei, Chen Li, Tianxin Huang, Yunsong Wang, and Gim Hee Lee. VCR-GauS: View consistent depth-normal regularizer for Gaussian surface reconstruction. In *Advances in Neural Information Processing Systems*, 2024. 3
- [10] Yihang Chen, Qianyi Wu, Weiyao Lin, Mehrtash Harber, and Jianfei Lu. HAC: Hash-grid assisted context for 3D Gaussian Splatting compression. In *European Conference on Computer Vision (ECCV)*, 2024. 4

- [11] Youyu Chen, Junjun Jiang, Kui Jiang, Xiao Tang, Zhihao Li, Xianming Liu, and Yinyu Nie. DashGaussian: Optimizing 3D Gaussian Splatting in 200 seconds. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 11146–11155, 2025. 3
- [12] Kai Cheng, Xiaoxiao Long, Kaizhi Yang, Yao Yao, Wei Yin, Yuexin Ma, Wenping Wang, and Xuejin Liu. GaussianPro: 3D Gaussian Splatting with progressive propagation. In *International Conference on Machine Learning (ICML)*, 2024. 3
- [13] Jaeyoung Chung, Jeongtaek Oh, and Kyoung Mu Lee. Depth-regularized optimization for 3D Gaussian Splatting in few-shot images. In *CVPR Workshop on 3D Vision and Modeling*, 2024. 3
- [14] Pinxuan Dai, Jiamin Xu, Wenxiang Xie, Xinguo Liu, Huamin Wang, and Weiwei Xu. High-quality surface reconstruction using Gaussian surfels. In *ACM SIGGRAPH 2024 Conference Proceedings*, 2024. 3
- [15] Xiaobin Deng, Changyu Diao, Min Li, Ruohan Yu, and Du-anqing Xu. Efficient density control for 3D Gaussian Splatting, 2024. 3
- [16] Zhiwen Fan, Kevin Wang, Kairun Wen, Zehao Zhu, De-jia Xu, and Zhangyang Wang. LightGaussian: Unbounded 3D Gaussian compression with 15x reduction and 200+ fps. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. 4
- [17] Guangchi Fang and Bing Wang. Mini-Splatting: Representing scenes with a constrained number of gaussians. In *European Conference on Computer Vision (ECCV)*, 2024. 3
- [18] Ben Fei et al. 3D Gaussian Splatting as a new era: A survey. *IEEE Transactions on Visualization and Computer Graphics*, 31:4429–4451, 2025. 2
- [19] Sara Fridovich-Keil, Alex Yu, Matthew Tancik, Qinhong Chen, Benjamin Recht, and Angjoo Kanazawa. Plenoxels: Radiance fields without neural networks. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022. 2, 3, 4
- [20] Jian Gao, Chun Gu, Youtian Lin, Hao Zhu, Xun Cao, Li Zhang, and Yao Yao. Relightable 3D gaussian: Real-time point cloud relighting with BRDF decomposition and ray tracing, 2023. 4
- [21] Sharath Girish, Kamal Bharadwaj, Abhinav Buddareddy-gari, Wen Shao, Deva Ramanan, and Ken Sakurada. EA-GLES: Efficient accelerated 3D Gaussians with lightweight encoding. In *European Conference on Computer Vision (ECCV)*, 2024. 4
- [22] Antoine Guédon and Vincent Lepetit. SuGaR: Surface-aligned Gaussian Splatting for efficient 3d mesh reconstruction and high-quality mesh rendering. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 2, 3
- [23] Alex Hanson, Allen Tu, Geng Lin, Vasu Singla, Matthias Zwicker, and Tom Goldstein. Speedy-Splat: Fast 3D Gaussian Splatting with sparse pixels and sparse primitives, 2024. 4
- [24] Alex Hanson, Allen Koressel, and Subhransu Maji. PUP 3D-GS: Principled uncertainty pruning for 3D Gaussian Splatting. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025. 3
- [25] Binbin Huang, Zehao Yu, Anpei Chen, Andreas Geiger, and Shenghua Gao. 2D Gaussian Splatting for geometrically accurate radiance fields. *ACM Transactions on Graphics (SIGGRAPH)*, 43(4), 2024. 2, 3, 7, 8
- [26] Yingwenqi Jiang, Jiadong Tu, Yuan Liu, Xifeng Gao, Xiaoxiao Long, Wenping Wang, and Yuexin Ma. GaussianShader: 3D gaussian splatting with shading functions for reflective surfaces. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 5322–5332, 2024. 4
- [27] Bernhard Kerbl, Georgios Kopanas, Thomas Leimkühler, and George Drettakis. 3D Gaussian Splatting for real-time radiance field rendering. *ACM Transactions on Graphics (SIGGRAPH)*, 42(4):1–14, 2023. 1, 2, 3, 4, 7
- [28] Shakiba Kheradmand, Daniel Rebain, Gopal Sharma, Weiwei Sun, Jeff Tseng, Hossam Isack, Abhishek Kar, Andrea Tagliasacchi, and Kwang Moo Yi. 3D Gaussian Splatting as Markov Chain Monte Carlo. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. 3
- [29] Joo Chan Lee, Daniel Rho, Xiangyu Sun, Jong Hwan Ko, and Eunbyung Park. Compact 3D Gaussian Representation for radiance field. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 4
- [30] Jiahe Li, Jiawei Zhang, Xiao Bai, Jin Zheng, Xin Ning, Jun Zhou, and Lin Gu. DNGaussian: Optimizing sparse-view 3D Gaussian radiance fields with global-local depth normalization. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. 3
- [31] Yanyan Li, Chenyu Lyu, Yan Di, Guangyao Zhai, Gim Hee Lee, and Federico Tombari. GeoGaussian: Geometry-aware Gaussian splatting for scene rendering. In *European Conference on Computer Vision (ECCV)*, 2024. 3
- [32] Zhihao Liang, Qi Zhang, Ying Feng, Ying Shan, and Kui Jia. Analytic-Splatting: Anti-aliased 3D Gaussian Splatting via analytic integration. In *European Conference on Computer Vision (ECCV)*, 2024. Oral. 2
- [33] Tao Lu, Mulin Yu, Linning Xu, Yuanbo Xiangli, Limin Wang, Dahua Lin, and Bo Dai. Scaffold-GS: Structured 3D Gaussians for view-adaptive rendering. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. Highlight. 2
- [34] Xiaoyang Lyu, Yang-Tian Sun, Yi-Hua Huang, Xiuzhe Wu, Ziyi Yang, Yilun Chen, Jiangmiao Pang, and Xiaojuan Qi. 3DGSr: Implicit surface reconstruction with 3D gaussian splatting. *ACM Transactions on Graphics*, 2024. 1, 2, 3, 4
- [35] Saswat Subhajyoti Mallick, Rahul Goel, Bernhard Kerbl, Markus Steinberger, Francisco Vicente Carrasco, and Fernando De La Torre. Taming 3DGS: High-quality radiance fields with limited resources. In *SIGGRAPH Asia 2024 Conference Papers*, 2024. 3
- [36] Ben Mildenhall, Pratul P Srinivasan, Matthew Tancik, Jonathan T Barron, Ravi Ramamoorthi, and Ren Ng. NeRF: Representing scenes as neural radiance fields for view synthesis. In *European Conference on Computer Vision (ECCV)*, pages 405–421. Springer, 2020. 2

- [37] Thomas Müller, Alex Evans, Christoph Schied, and Alexander Keller. Instant neural graphics primitives with a multiresolution hash encoding. pages 1–15, 2022. [3](#), [4](#)
- [38] Seungtae Nam, Xiangyu Sun, Gyeongjin Kang, Younggeun Lee, Seungjun Oh, and Eunbyung Park. Generative densification: Learning to densify Gaussians for high-fidelity generalizable 3d reconstruction, 2024. [3](#)
- [39] Lukas Radl, Michael Steber, Mathias Parger, Bernhard Kerbl, and Markus Steinberger. StopThePop: Sorted Gaussian Splatting for view-consistent real-time rendering. *ACM Transactions on Graphics*, 43(4), 2024. [2](#)
- [40] Kerui Ren, Lihan Jiang, Tao Lu, Mulin Yu, Linning Xu, Zhangkai Ni, and Bo Dai. Octree-GS: Towards consistent real-time rendering with LOD-structured 3D Gaussians. *IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI)*, 2025. [2](#)
- [41] Matthew Tancik, Ethan Weber, Evonne Ng, Ruilong Li, Brent Yi, Terrance Wang, Alexander Kristoffersen, Jake Austin, Kamyar Salahi, Abhik Ahuja, et al. Nerfstudio: A modular framework for neural radiance field development. In *ACM SIGGRAPH*, 2023. [2](#), [6](#), [7](#)
- [42] Peng Wang, Lingjie Liu, Yuan Liu, Christian Theobalt, Taku Komura, and Wenping Wang. NeuS: Learning neural implicit surfaces by volume rendering for multi-view reconstruction. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021. [2](#), [3](#), [4](#)
- [43] Penguin Wang, Yueji Wang, Dongjin Wang, Shijia Mohan, Zhiwen Fan, Liang Wu, Ruijin Cai, Yung-Yu Yeh, Zhangyang Wang, Qixing Liu, and Rakesh Ranjan. Steepest descent density control for compact 3D Gaussian Splatting. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025. [3](#)
- [44] Yiwen Wang, Zilong Chen, Liang Chi, Zhongang Wang, Ziwei Liu, and Guosheng Yang. GaussianEditor: Swift and controllable 3d editing with Gaussian Splatting. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. [4](#)
- [45] Guikun Wu and Jie Yi. A survey on 3D Gaussian Splatting. *arXiv preprint arXiv:2401.03890*, 2024. [2](#)
- [46] Tianyi Xie, Zeshun Zong, Yuxing Qiu, Xuan Li, Yutao Feng, Yin Yang, and Chenfanfu Jiang. PhysGaussian: Physics-integrated 3D gaussians for generative dynamics. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. Highlight. [4](#)
- [47] Vickie Ye, Matias Turkulainen, and the nerfstudio team. gsplat: An open-source library for Gaussian Splatting. *Journal of Machine Learning Research (JMLR)*, 26:1–15, 2025. [2](#), [4](#), [6](#), [7](#)
- [48] Zongxin Ye, Wenyu Li, Sidun Liu, Peng Qiao, and Yong Dou. AbsGS: Recovering fine details for 3D Gaussian Splatting. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. [2](#), [3](#)
- [49] Mulin Yu, Tao Lu, Linning Xu, Lihan Jiang, Yuanbo Xiangli, and Bo Dai. GSDF: 3DGS meets SDF for improved rendering and reconstruction. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. [1](#), [2](#), [3](#), [4](#)
- [50] Zehao Yu, Anpei Chen, Binbin Huang, Torsten Sattler, and Andreas Geiger. Mip-Splatting: Alias-free 3D Gaussian Splatting. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. Best Student Paper. [2](#)
- [51] Zehao Yu, Torsten Sattler, and Andreas Geiger. Gaussian Opacity Fields: Efficient adaptive surface reconstruction in unbounded scenes. *ACM Transactions on Graphics*, 43(6), 2024. [2](#), [3](#)
- [52] Baowen Zhang, Chuan Fang, Rakesh Shrestha, Yixun Liang, Xiaoxiao Long, and Ping Tan. RaDe-GS: Rasterizing depth in Gaussian splatting. *ACM Transactions on Graphics*, 2024. [3](#)
- [53] Jiawei Zhang, Jiahe Li, Xiaohan Yu, Lei Huang, Lin Gu, Jin Zheng, and Xiao Bai. CoR-GS: Sparse-view 3D Gaussian Splatting via co-regularization. In *European Conference on Computer Vision (ECCV)*, 2024. [3](#)
- [54] Jiahui Zhang, Fangneng Zhan, Muyu Xu, Shijian Lu, and Eric Xing. FreGS: 3D Gaussian Splatting with progressive frequency regularization. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. [3](#)
- [55] Wenyuan Zhang, Yu-Shen Liu, and Zhizhong Han. Neural Signed Distance Function Inference through Splatting 3D Gaussians Pulled on Zero-Level Set. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. [1](#), [2](#), [3](#), [4](#), [9](#), [10](#)
- [56] Zheng Zhang, Wenbo Hu, Yixing Lao, Tong He, and Hengshuang Zhao. Pixel-GS: Density control with pixel-aware gradient for 3D Gaussian Splatting. In *European Conference on Computer Vision (ECCV)*, 2024. [2](#), [3](#)
- [57] Zehao Zhu, Zhiwen Fan, Yifan Jiang, and Zhangyang Wang. FSGS: Real-time few-shot view synthesis using Gaussian Splatting. In *European Conference on Computer Vision (ECCV)*, 2024. [3](#)
- [58] Matthias Zwicker, Hanspeter Pfister, Jeroen Van Baar, and Markus Gross. EWA splatting. *IEEE Transactions on Visualization and Computer Graphics*, 8(3):223–238, 2002. [2](#), [4](#)